



TU Dublin Operating Systems COMP 3602

Lecture 5 Threads





Process (definitions)

- Process is a program in execution (the unit of work in a modern computing system)
 - A program is a file stored on a hard disk and containing a list of instructions (i.e. executable file)
 - A program becomes a process when an executable file is loaded into memory.





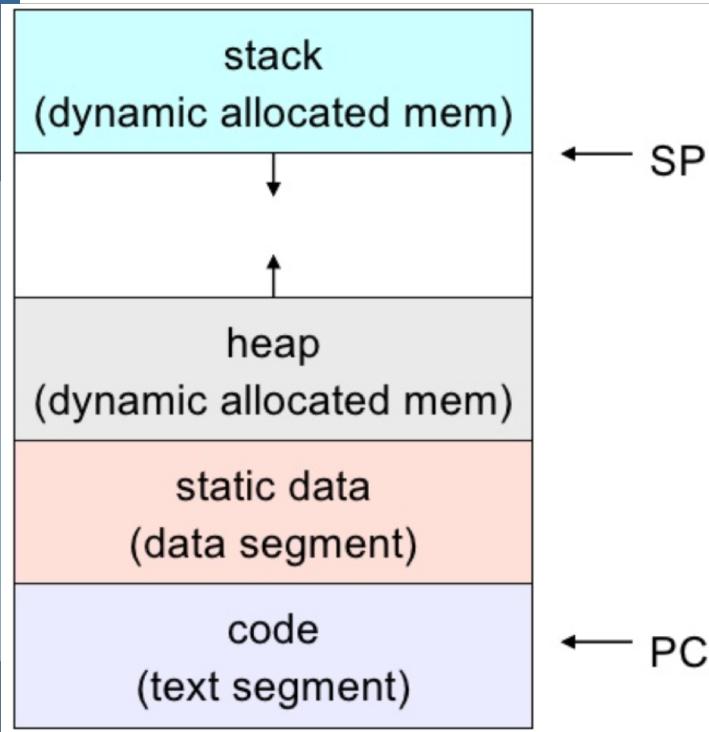
Process components

- A process consists of everything needed to run the program
 - An address space (sections in memory)
 - ▶ the data needed for the running program
 - heap
 - stack
 - static data
 - ▶ the code (instructions) for the running program
 - CPU state (values of CPU registers)
 - ▶ The program counter (PC), indicating the next instruction
 - ▶ The stack pointer (SP)
 - ▶ Other general purpose register
 - A set of OS resources
 - ▶ open files
 - ▶ I/O devices





The memory layout of a process

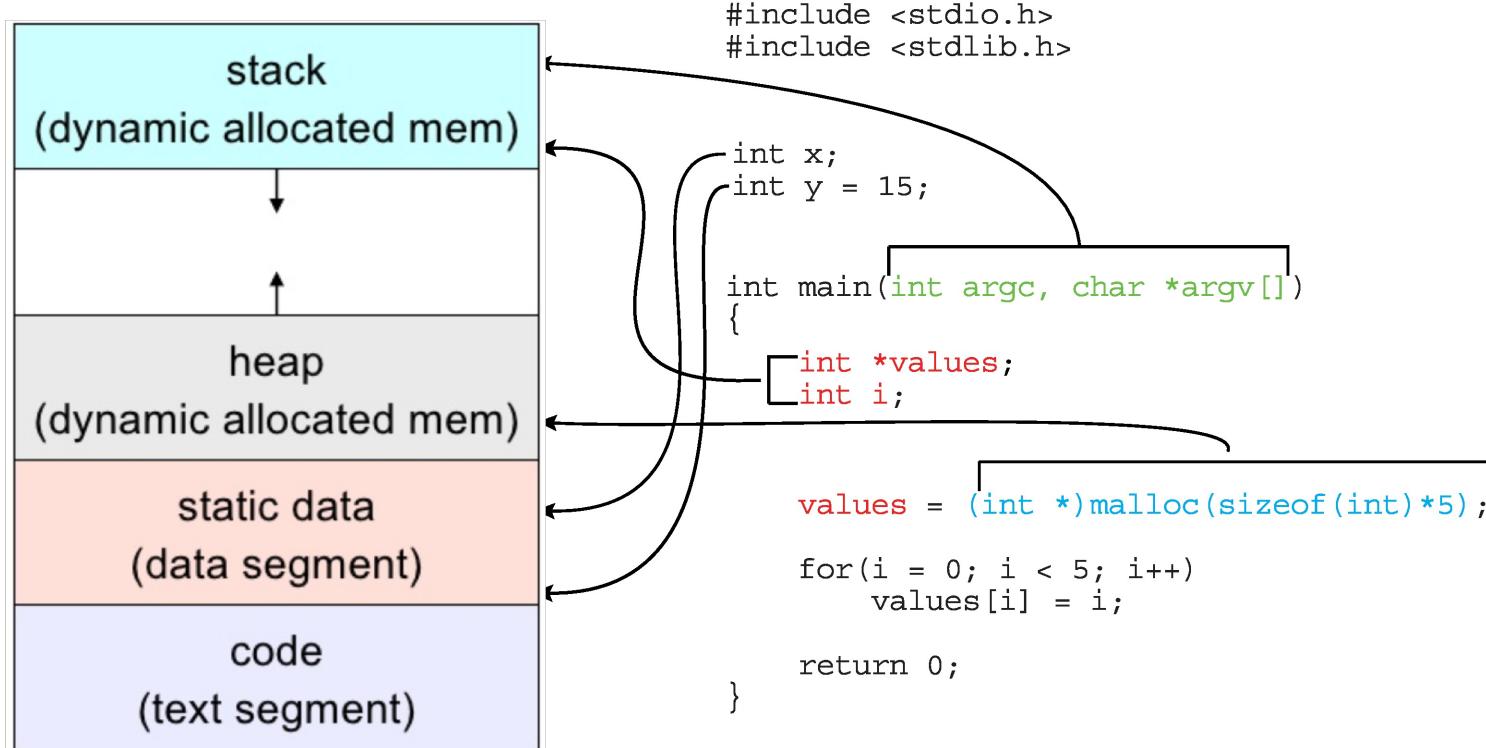


- Memory sections needed for the running program
 - **stack**—temporary data (local variables, function parameters, return addresses)
 - **heap**—memory that is dynamically allocated during program run time
 - **static data**—global variables
- **Text section**—the code (instructions) for the running program
- **CPU registers:**
 - ▶ The program counter (PC), indicating the next instruction
 - ▶ The stack pointer (SP)





Memory Layout of a C Program

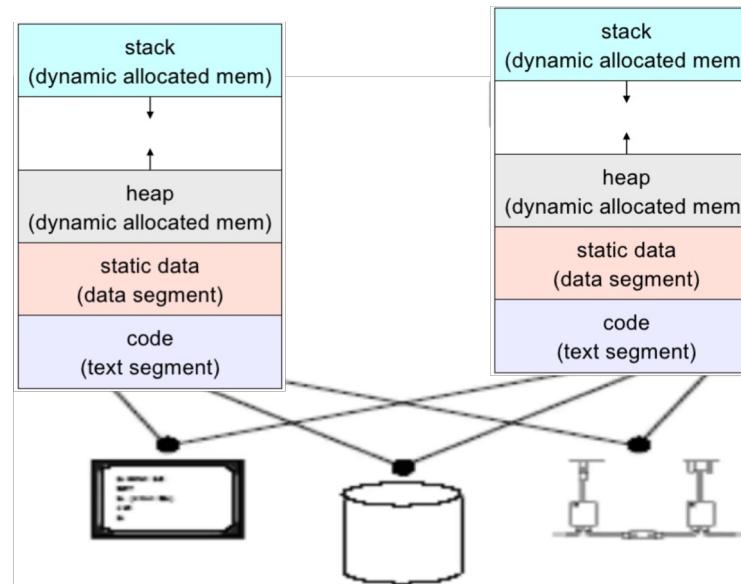




Process creation using *fork()*

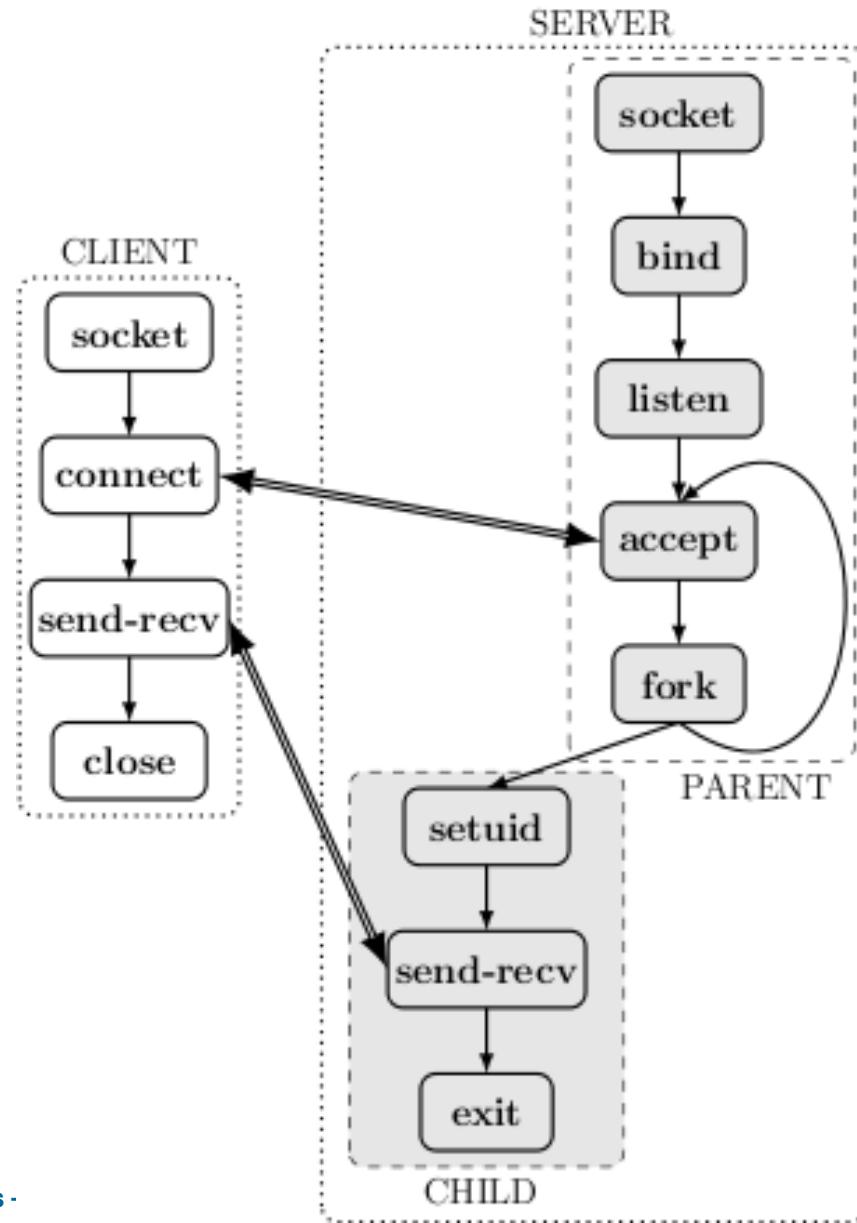
- *fork()* system call allows one process, *the parent*, to create a new process, *the child*
- *the child process* is a copy of the address space of *the parent* (code, static data, heap, stack)

```
pid = fork();
if (pid < 0) {
    //error
}
else if (pid == 0) {
    // child process
}
else {
    // parent process
}
```



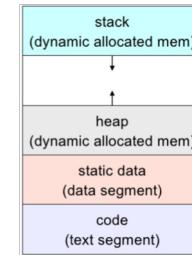
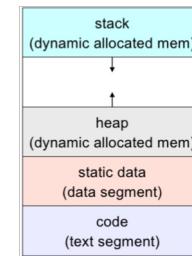
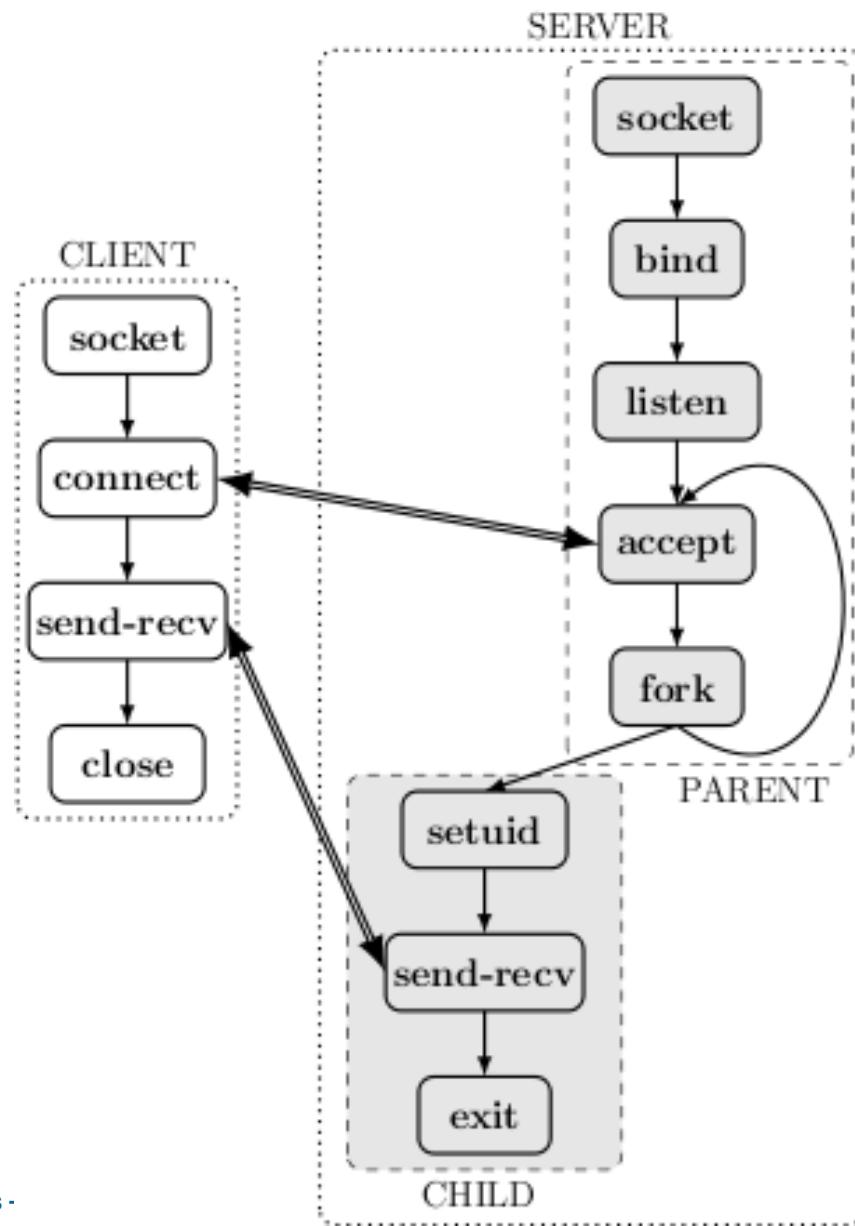


Web server with fork() system call



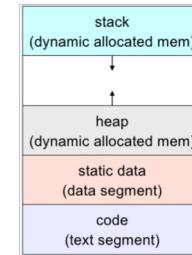
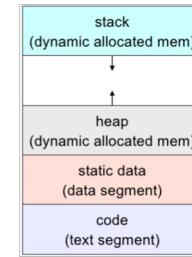
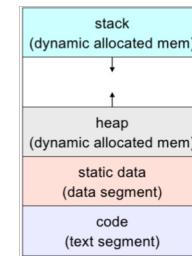
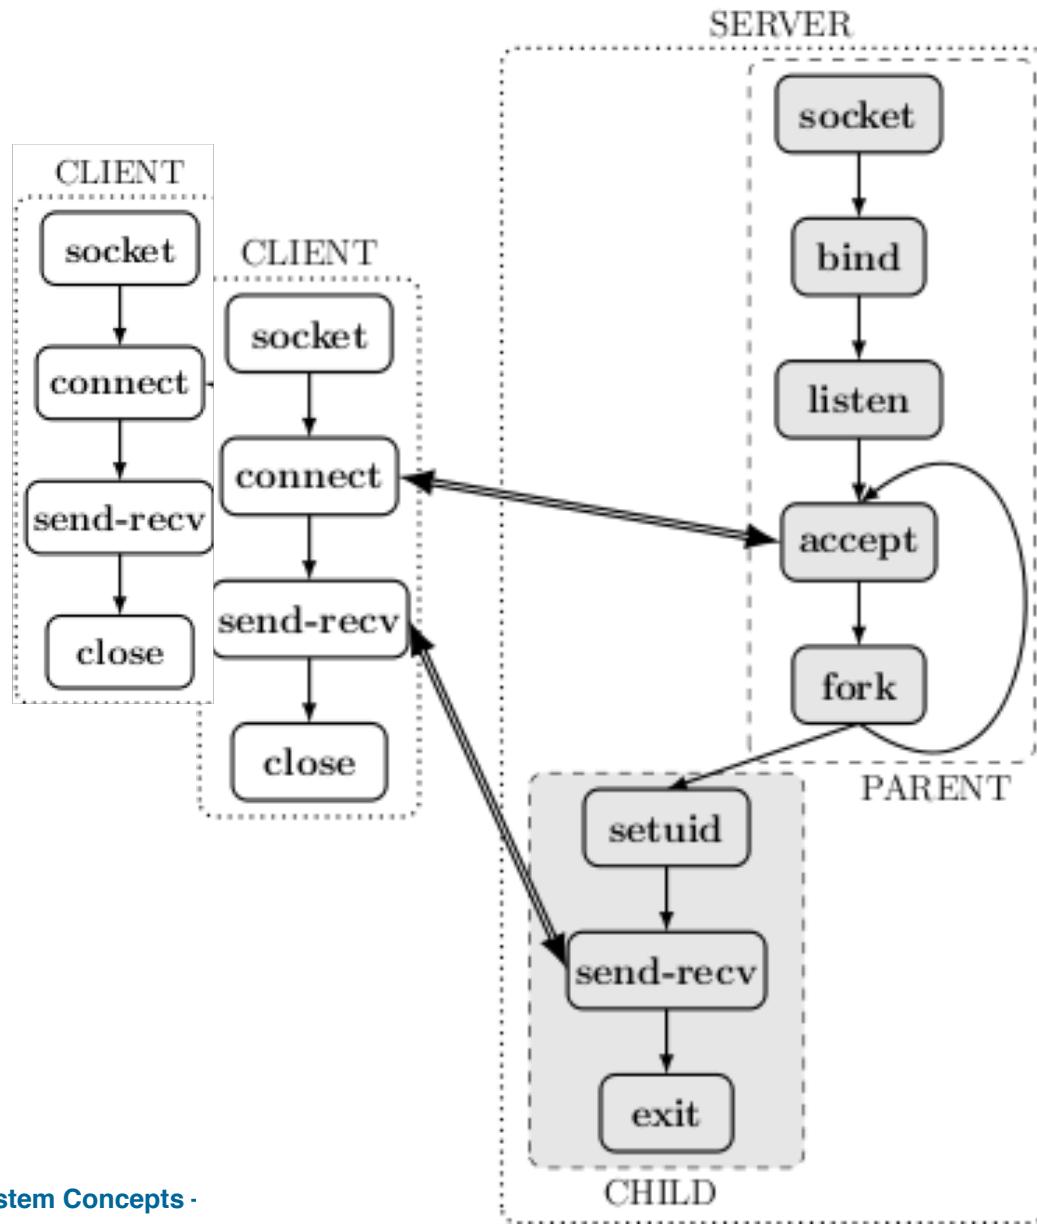


Web server with fork() system call



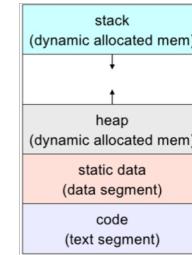
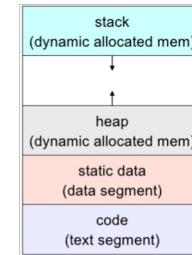
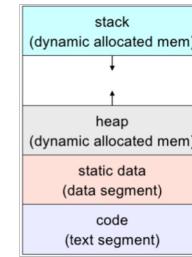
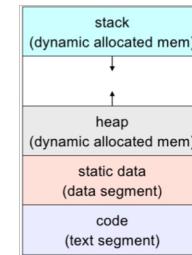
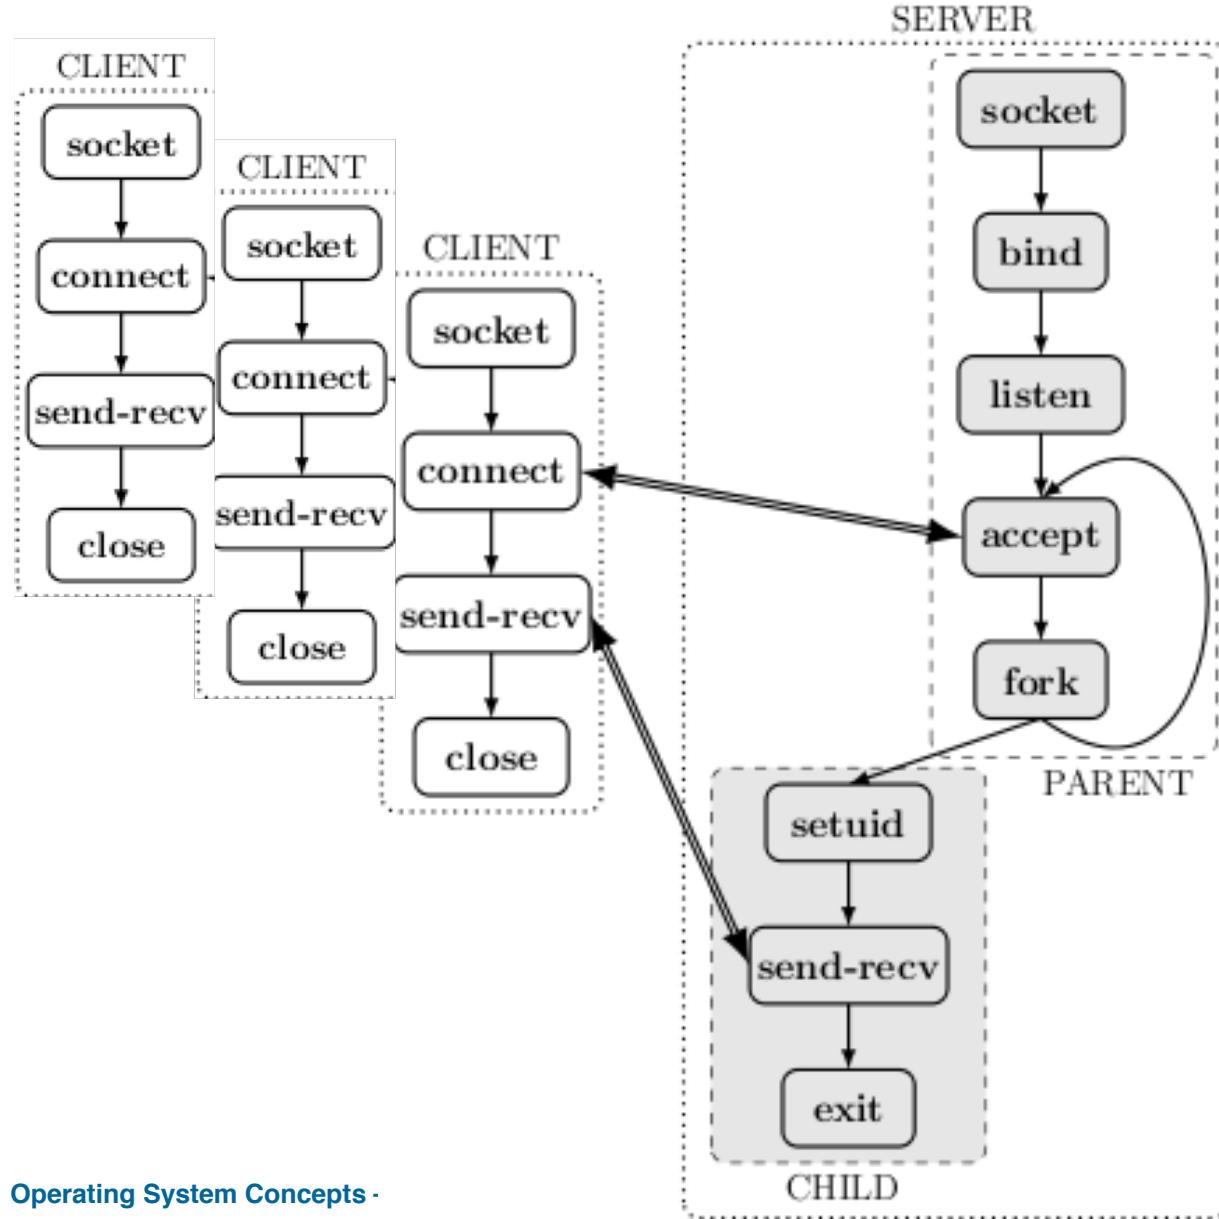


Web server with fork() system call





Web server with fork() system call





Processes overhead

- Imagine a web server handling 1000 requests at the same time
 - 1000 child processes
 - All processes run the same code
 - All processes access the same data
 - All processes use the same resources

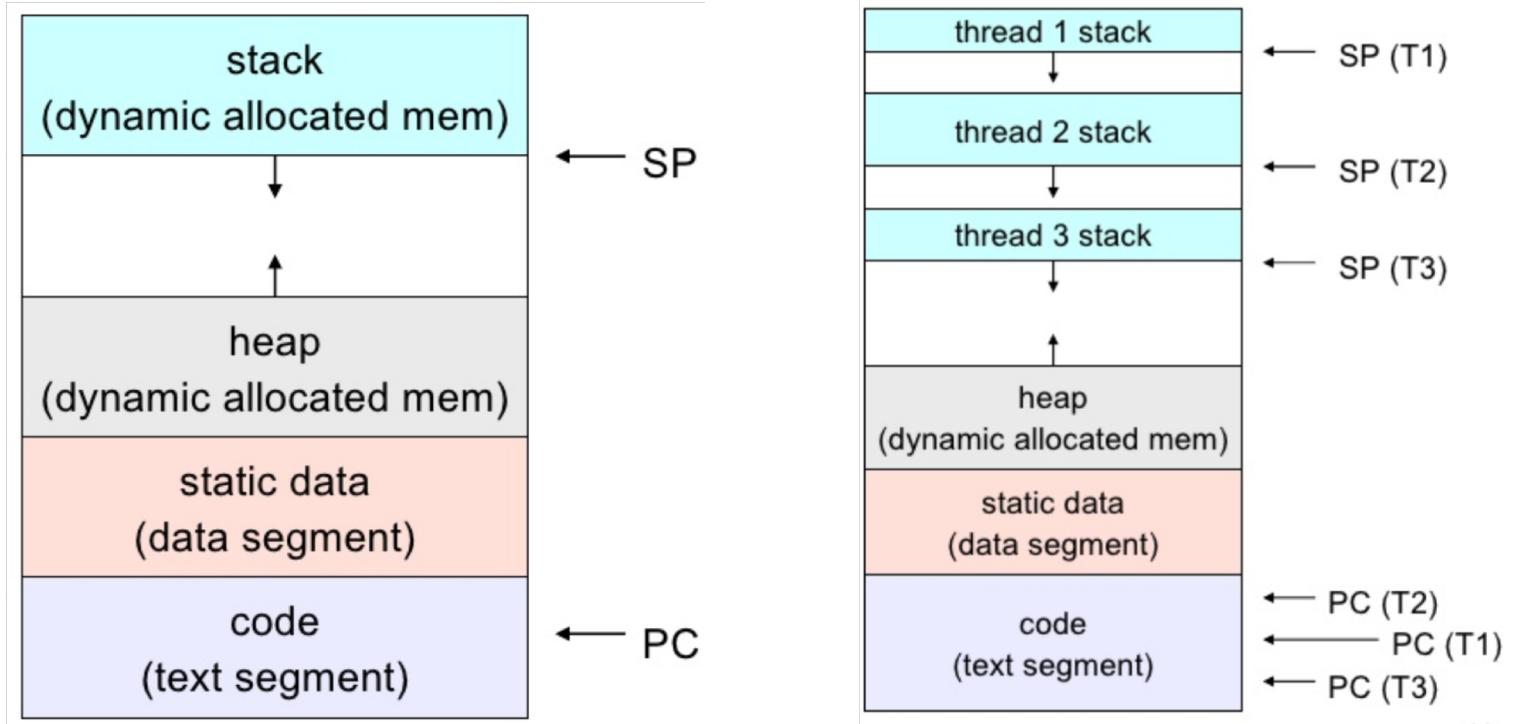
- What's the point in having 1000 child processes? Inefficient use of resources!
 - process creation is heavy-weight
 - memory:
 - ▶ 1000 copies of the same address space,
 - ▶ 1000 PCBs (process control blocks)
 - time:
 - ▶ copying 1000 times the same address space is time consuming!





Threads

- A **thread** is the smallest unit of processing that can be performed in an OS (a basic unit of CPU utilization).
- A traditional process performs one task at a time, i.e. has a single thread of control.
- If a process has multiple threads of control, it can perform more than one task at a time.





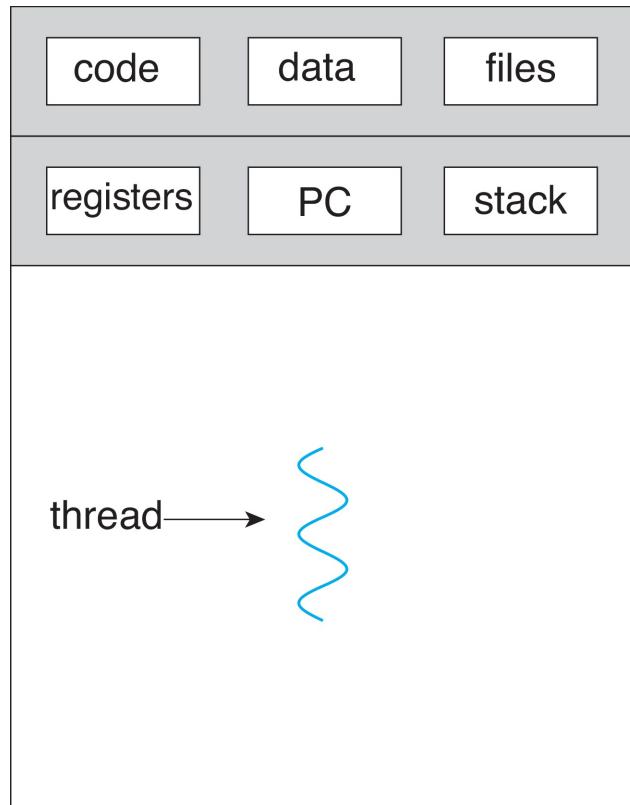
Thread components

- A thread is a flow of execution through the process code with its own:
 - thread ID
 - program counter (PC)
 - stack
 - register set
- Each thread belongs to exactly one process; cannot exist outside a process
- A thread shares with other threads belonging to the same process:
 - its code section
 - data section
 - operating-system resources
 - ▶ open files and signals
 - ▶ I/O devices

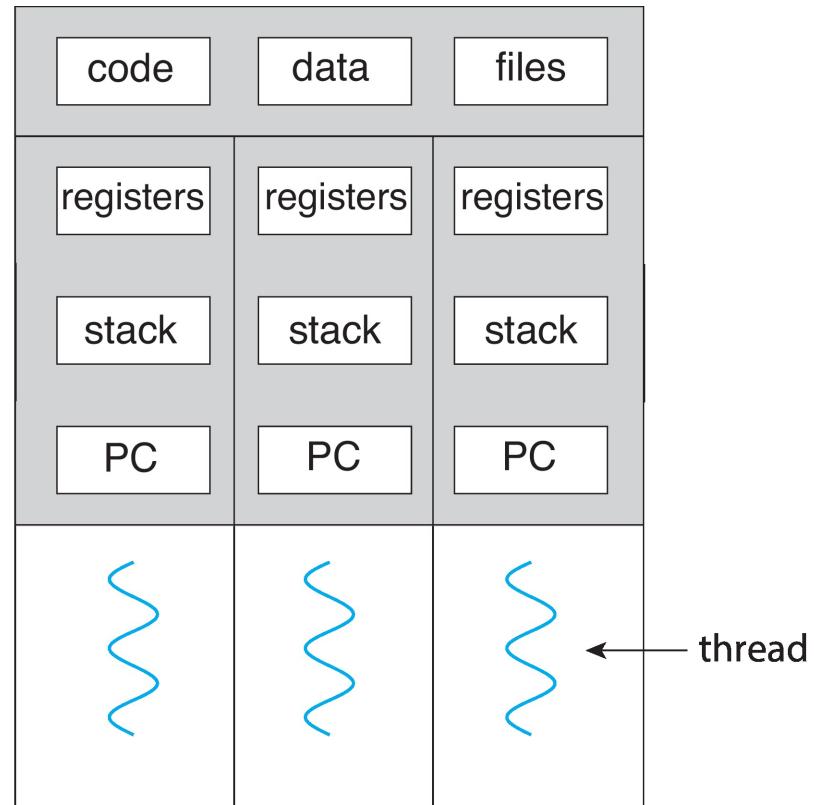




Single and Multithreaded Processes



single-threaded process



multithreaded process





Processes and Threads

- Most modern OS's support both:
 - processes (which define the address space)
 - threads (which define a sequential execution stream within a process)
- Process creation is heavy-weight while thread creation is light-weight
- Process switching needs interaction with OS while thread switching doesn't.
- A thread runs in a single process (address space)
 - One process can have multiple threads and they can read, write, change another thread's data.
- All threads in one process share system resources
 - creating threads is quick (ca. 10 times quicker than processes)
 - ending threads is quick
 - switching threads within one process is quick
 - inter-thread communication is quick and easy (shared memory)





Benefits

■ Economy

- thread creation is cheaper (time, memory) than process creation
- thread switching has lower overhead than context switching
- better utilization of multicore (multiprocessor) architecture

■ Resource Sharing

- threads share the same address space
- no need to implement Inter-Process Communications

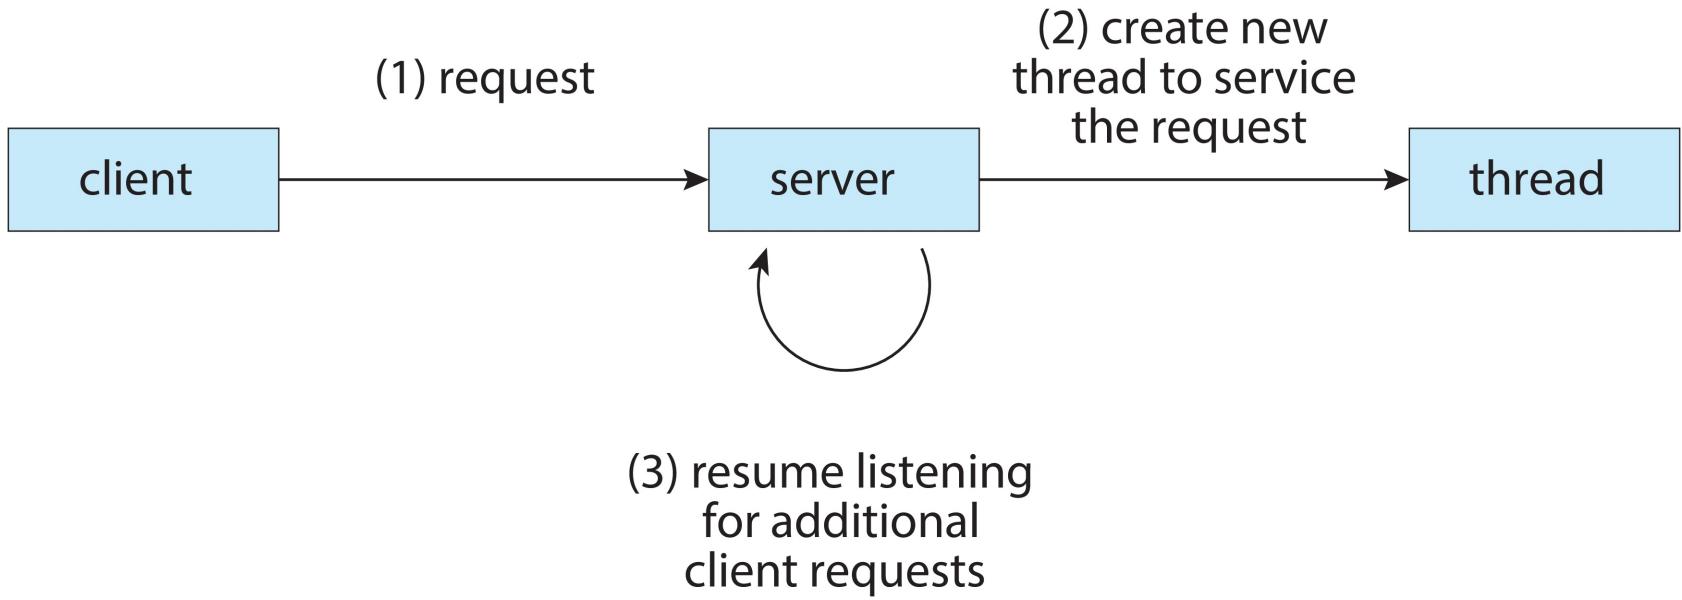
■ Responsiveness – interactive applications can be performing two tasks at the same time (important for user interfaces)

■ Scalability – multithreading can take advantage of multicore architectures





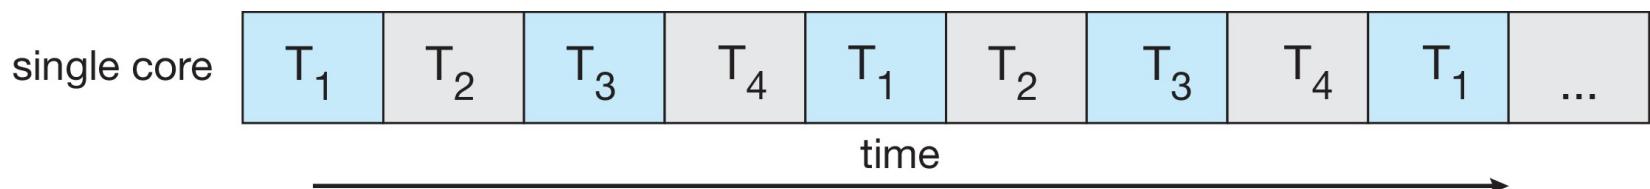
Multithreaded Server Architecture



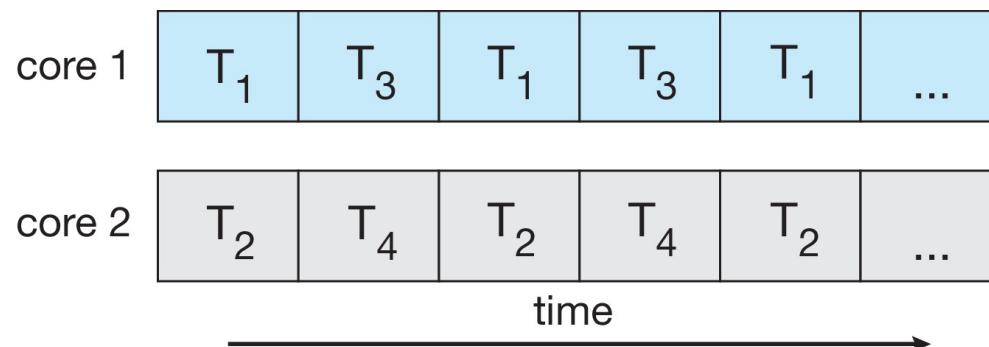


Concurrency vs. Parallelism

- **Concurrent execution on single-core system** (more than one task can make progress when a task scheduler provides concurrency)



- **Parallelism on a multi-core system** - a system can perform more than one task simultaneously

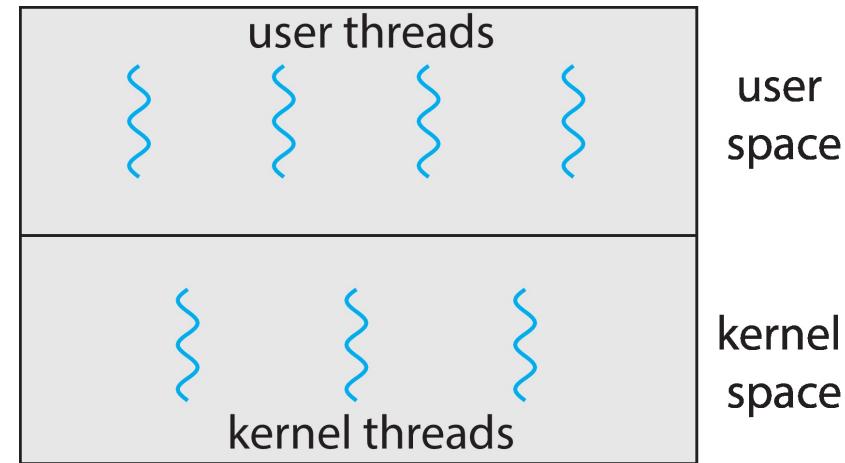




Thread types

■ **User threads** – managed by user-level threads library

- Three primary thread libraries:
 - ▶ POSIX **Pthreads**
 - ▶ Windows threads
 - ▶ Java threads



■ **Kernel threads** - supported by the Kernel

- Examples:
 - ▶ Windows
 - ▶ Linux
 - ▶ Mac OS X
 - ▶ iOS
 - ▶ Android





Multithreading Models

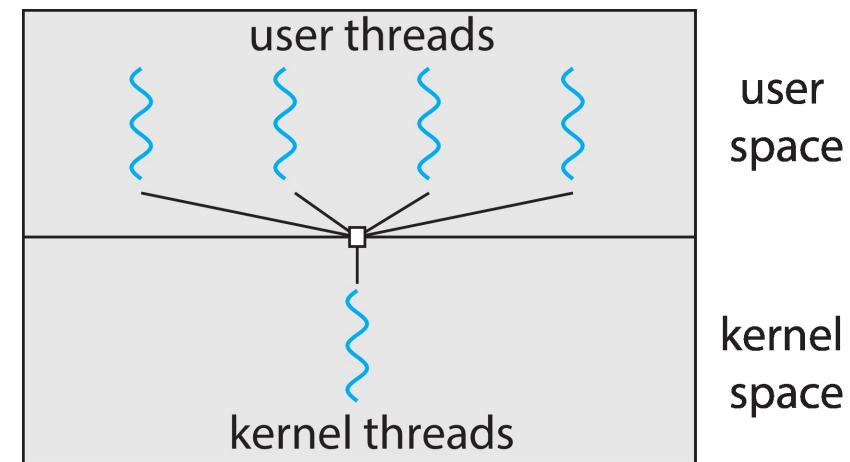
- Operating Systems provide various types of mapping between user-level threads and kernel-level threads:
 - Many-to-One
 - One-to-One
 - Many-to-Many





Many-to-One

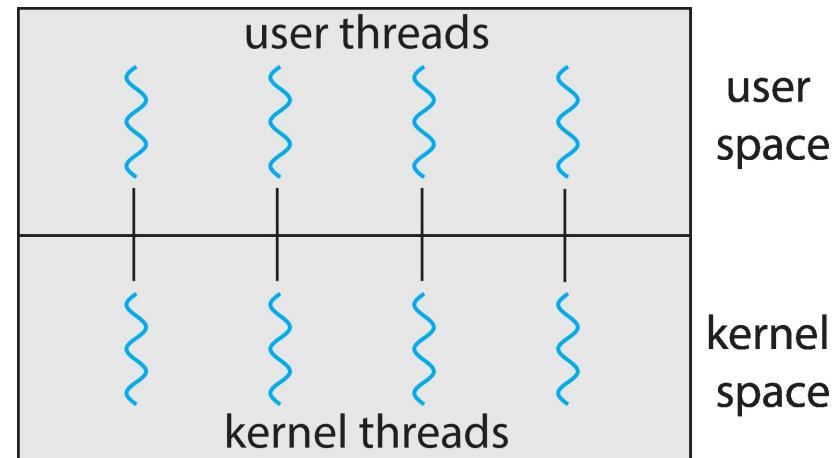
- Many user-level threads mapped to single kernel thread
 - If single kernel thread blocks all the other cannot run
- Multiple threads cannot run in parallel on multicore systems because of single kernel thread
- Few systems currently use this model
 - **Solaris Green Threads**
 - **GNU Portable Threads**





One-to-One

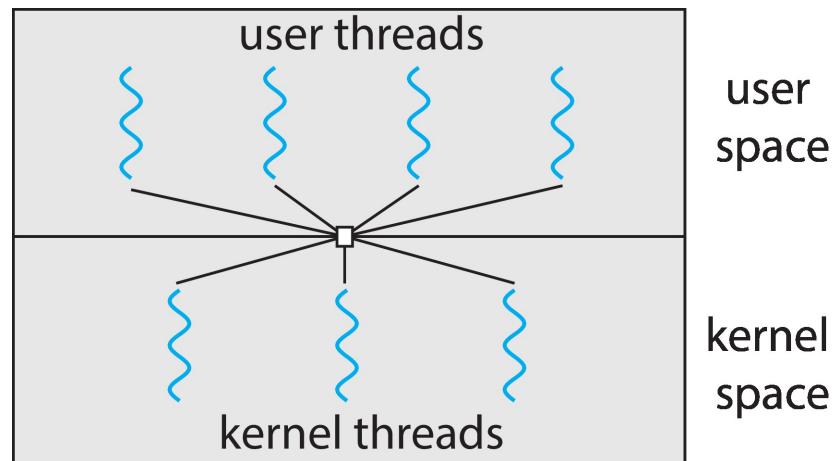
- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
 - Windows
 - Linux





Many-to-Many Model

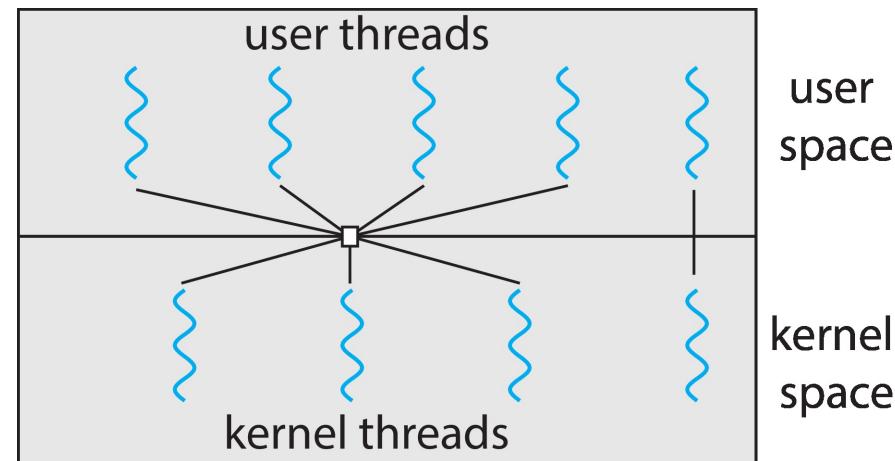
- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads





Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- Examples
 - IRIX
 - Tru64 UNIX
 - Solaris 8 and earlier





Thread Implementation

- Two primary ways of implementing
 - Threads can be implemented as part of the OS
 - ▶ Kernel-level library supported by the OS
 - Threads can be implemented in user space
 - ▶ user-space libraries
-





PThreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Linux & Mac OS X)





Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid,NULL);

    printf("sum = %d\n",sum);
}
```





Pthreads Example (cont)

```
/* The thread will execute in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```





Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```



THREAD PROGRAMMING

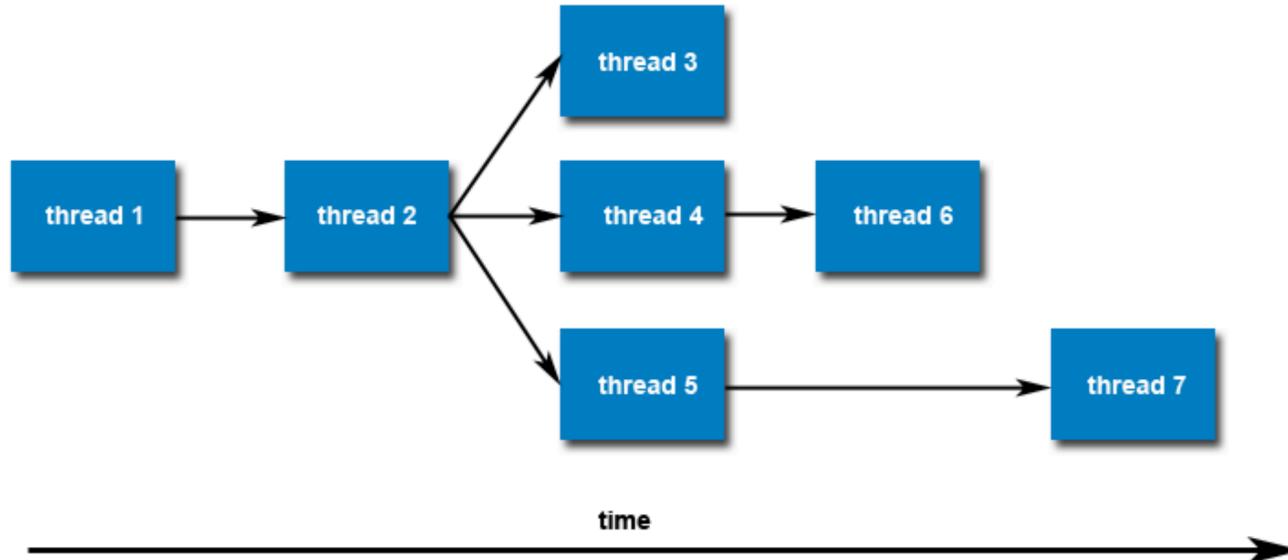
POSIX Threads (Pthreads) API

- Portable Operating System Interface for UNIX
 - POSIX for thread programming, specified by the IEEE POSIX 1003.1c standard
 - C Language
 - Supported by most operating systems: Linux, Mac OS X, Windows (partially), etc.
- The Pthreads API contains over 60 subroutines in three major classes:
 - Thread management: creating, terminating, joining
 - Mutexes: provides exclusive access to code segments/variables with the use of locks (mutual exclusion)
 - Condition variables: provides synchronization and communication between threads that share a mutex

Creating threads

`pthread_create (thread_id, attr, function, arg)`

- Creates a child thread which starts running the specified **function**.
- Once created, threads may create other child threads.



Creating threads (cont.)

```
int pthread_create(pthread_t *thread_id, pthread_attr_t *attr,  
                  void *(*function)(void *), void *arg);
```

- Function ***pthread_create()*** returns “0” when successful or error integer number otherwise
- Arguments:
 - ***pthread_t *thread***: address of a threadID
 - ***pthread_attr_t *attr***: attributes that can be applied to this thread (can be NULL)
 - ***void *(*function)(void *)***: the actual function this thread executes
 - ***void *arg***: pointer to arguments to be passed to the actual function (can be NULL)

Terminating threads

`void pthread_exit(void *status);`

- Function **`pthread_exit()`** terminates the thread
- The **`status`** value is returned to the master thread if the master thread is waiting (by calling **`pthread_join()`**) for thread termination
- Ways for a thread to terminate:
 - The thread function calls the **`return`** statement
 - The main() process calls **`return`** or **`exit()`**
 - The thread function calls **`pthread_exit()`**
 - The thread is canceled by another thread via the **`pthread_cancel()`** function

Compiling Pthreads:

```
$ gcc <source_code.c> -o <executable> -lpthread
```

Simplest thread code

```
#include<pthread.h>
#include<stdio.h>

void * function(void * arg){
    printf("Hello world, this is child thread!\n");

    //terminate thread
    printf("Child thread is exiting\n");
    pthread_exit(NULL);
}

int main(){
    printf("Hello world, this is master thread!\n");
    pthread_t t_id;

    //create thread
    int r = pthread_create(&t_id, NULL, function, NULL);
    if (r == 0)
        printf("Child thread created, t_id = %d\n", (int)t_id);
    else
        printf("Thread error nr: %d\n", r);

    printf("Master thread is exiting\n");
    return 0;
}
```

OUTPUT, version 1

Hello world, this is master thread!
Child thread created, t_id = 174768128
Hello world, this is child thread!
Child thread is exiting
Master thread is exiting

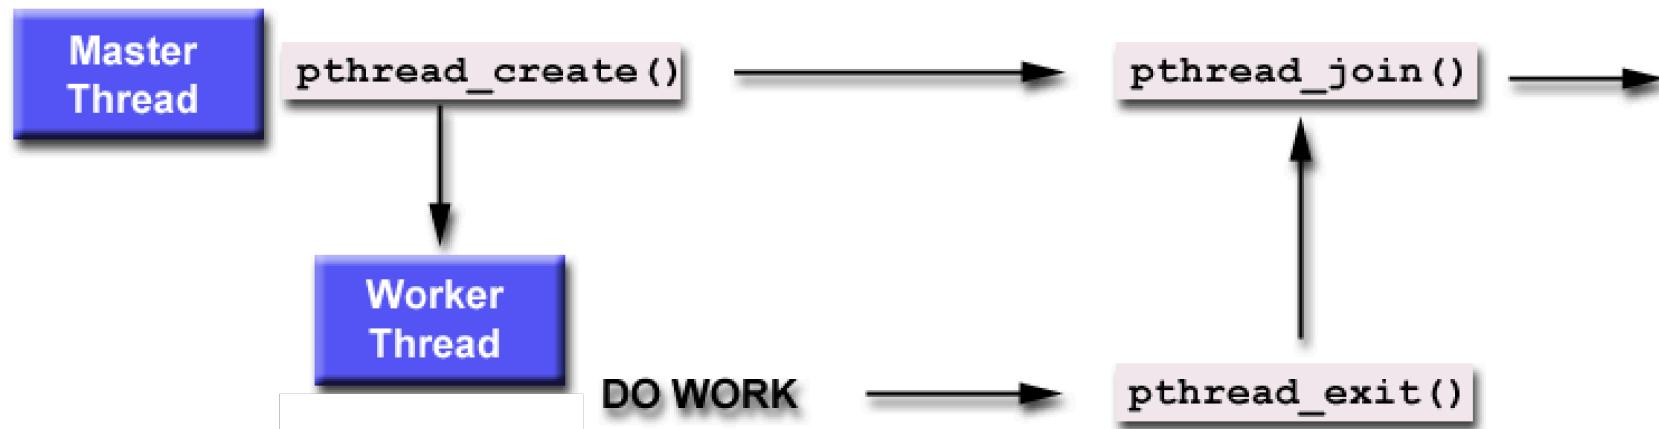
OUTPUT, version 2

Hello world, this is master thread!
Child thread created, t_id = 73080832
Master thread is exiting

Blocking the master thread until the child thread terminates.

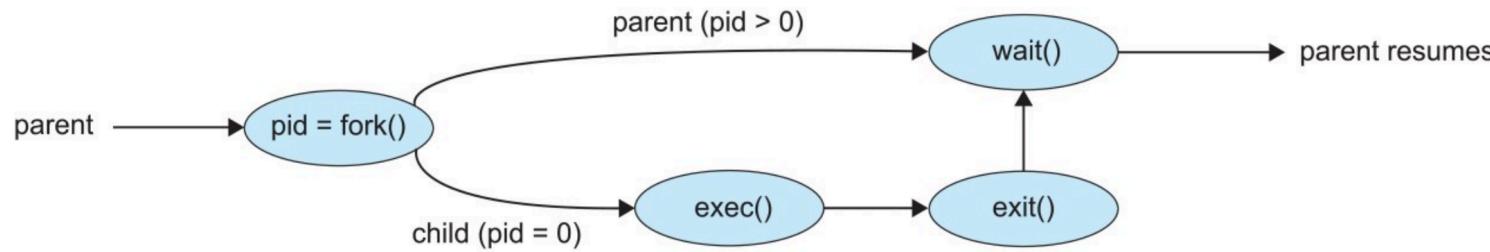
```
int pthread_join(pthread_t thread, void **status);
```

- blocks the calling (master) thread
- the calling thread waits for successful termination of the thread specified as the first argument
- an optional *status value is received from the terminating thread's call

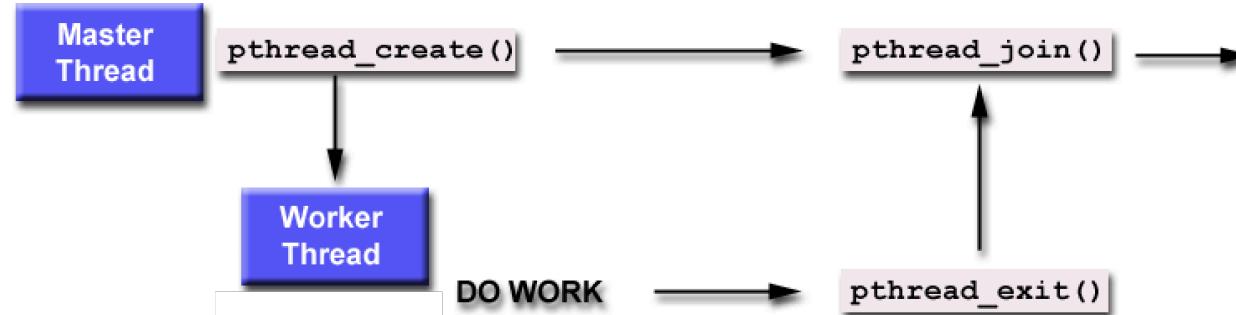


wait() system call & *pthread_join()* function

- The parent process may wait for termination of a child process by using the ***wait()*** system call.



- The parent (master) thread may wait for termination of a child thread by using the ***pthread_join()*** function.



Simplest thread code with *pthread_join()*

```
#include<pthread.h>
#include<stdio.h>

void * function(void * arg){
    printf("Hello world, this is child thread!\n");
    //terminate thread
    printf("Child thread is exiting\n");
    pthread_exit(NULL);
}

int main(){
    printf("Hello world, this is master thread!\n");
    pthread_t t_id;

    //create thread
    int r = pthread_create(&t_id, NULL, function, NULL);
    if (r == 0)
        printf("Child thread created, t_id = %d\n", (int)t_id);
    else
        printf("Thread error nr: %d\n", r);

    pthread_join(t_id, NULL);
    printf("Master thread is exiting\n");
    return 0;
}
```

OUTPUT

```
Hello world, this is master thread!
Child thread created, t_id = 149229568
Hello world, this is child thread!
Child thread is exiting
Master thread is exiting
```

Multi-threaded process (2 threads)

```
#include<pthread.h>
#include<stdio.h>

void * f1(void * arg){
    for (int i=0; i<1000; i++){
        printf("1");
        fflush(stdout);

    }
    pthread_exit(NULL);
}

void * f2(void * arg){
    for (int i=0; i<1000; i++){
        printf("2");
        fflush(stdout);

    }
    pthread_exit(NULL);
}

int main(){
    pthread_t t_id1, t_id2;
    int r1 = pthread_create(&t_id1, NULL, f1, NULL);
    int r2 = pthread_create(&t_id2, NULL, f2, NULL);
    pthread_join(t_id1, NULL);
    pthread_join(t_id2, NULL);
return 0;
}
```

OUTPUT

22221111222
22222222222222222222222212111111111111111111
1111111111211111112111122222111111122222
2222222222222222222211111111111111111111111
22222222222222222222211222111122222
211111111111111122222222222212122222222
222221112111111122222222222222222222
2222211111111111111111111111111111111111111
12211111111122222222222222212222222
22222222222222211111111111111111222222
12221222211222222222222222222221111111
1111111111211111211111222222222222211
1111111111111111111111222222222222222
22222222222222222222222222221111112222
111122221111111111111122222222222221
11222222222222211111112111111111111111

Multi-threaded process (2 threads)

```
#include<pthread.h>
#include<stdio.h>
#include <unistd.h>

void * f1(void * arg){
    for (int i=0; i<10; i++){
        printf("1");
        fflush(stdout);
        sleep(1)
    }
    pthread_exit(NULL);
}

void * f2(void * arg){
    for (int i=0; i<10; i++){
        printf("2");
        fflush(stdout);
        sleep(1);
    }
    pthread_exit(NULL);
}

int main(){
    pthread_t t_id1, t_id2;
    int r1 = pthread_create(&t_id1, NULL, f1, NULL);
    int r2 = pthread_create(&t_id2, NULL, f2, NULL);
    pthread_join(t_id1, NULL);
    pthread_join(t_id2, NULL);
return 0;
}
```

OUTPUT

```
$ time ./thread4
12212121212112121221
real 0m10.034s
user 0m0.002s
sys 0m0.005s
```

Single-threaded process (2 threads)

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>

void f1(){
    for (int i=0; i<10; i++){
        printf("1");
        fflush(stdout);
        sleep(1);
    }
}

void f2(){
    for (int i=0; i<10; i++){
        printf("2");
        fflush(stdout);
        sleep(1);
    }
}

int main(){
    f1();
    f2();
    return 0;
}
```

OUTPUT

```
$ time ./thread5
11111111112222222222
real 0m20.053s
user 0m0.002s
sys 0m0.004s
```

Passing arguments to the thread function

- Passing an integer

```
int myInteger = 10;  
int * ptr = &myInteger;  
pthread_create(&tid, NULL, fn, (void *) ptr);
```

- Reading an integer in the thread

```
void *fn (void *ptr){  
    int * x = (int *)ptr;  
    printf("I received the value: %d", *x);  
    (...)
```

- Passing a string

```
char message[25] = {"hello mr thread"};  
pthread_create(&t_id, NULL, fn, message);
```

- Reading string in the thread

```
void *function (void *ptr){  
    printf("I received a message: %s", ptr);  
    (...)
```

Thread code with integer passing

```
#include<pthread.h>
#include<stdio.h>

void * function(void * ptr){
    printf("Hello world, this is child thread!\n");
    int *x = (int*) ptr;
    printf("I received an integer: %d\n", *x);
    pthread_exit(NULL);
}

int main(){
    printf("Hello world, this is master thread!\n");
    pthread_t t_id;

    //a variable x stores the integer value
    int x = 10;

    // a pointer to the address of x will be passed
    int * ptr = &x;

    int r = pthread_create(&t_id, NULL, function, (void *)ptr);
    pthread_join(t_id, NULL);
    printf("Master thread is exiting\n");
return 0;
}
```

OUTPUT

Hello world, this is master thread!
Hello world, this is child thread!
I received an integer: 10
Master thread is exiting

Thread code with string passing

```
#include<pthread.h>
#include<stdio.h>

void * function(void * ptr){
    printf("Hello world, this is child thread!\n");

    printf("I got a message: %s\n", ptr);
    pthread_exit(NULL);
}

int main(){
    printf("Hello world, this is master thread!\n");
    pthread_t t_id;

    //a string "hello mr thread" will be sent
    char message[25] = {"hello mr thread"};

    //create thread
    int r = pthread_create(&t_id, NULL, function, message);

    pthread_join(t_id, NULL);
    printf("Master thread is exiting\n");
return 0;
}
```

OUTPUT

Hello world, this is master thread!
Hello world, this is child thread!
I got a message: hello mr thread
Master thread is exiting

Multi-threaded process with integer passing

```
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <pthread.h>
#include <unistd.h>

#define NUM_THREADS 5

void *perform_work(void *arguments) {
    int index = *((int *)arguments);
    int sleep_time = 1 + rand() % NUM_THREADS;
    printf("THREAD %d: Started.\n", index);
    printf("THREAD %d: Will be sleeping for %d seconds.\n", index, sleep_time);
    sleep(sleep_time);
    printf("THREAD %d: Ended.\n", index);
}
```

Multi-threaded process with integer passing (cont.)

```
int main(void) {
    pthread_t threads[NUM_THREADS];
    int thread_args[NUM_THREADS], i, result_code;

//create all threads one by one
    for (i = 0; i < NUM_THREADS; i++) {
        printf("IN MAIN: Creating thread %d.\n", i);
        threads[i] = i;
        thread_args[i] = i;
        result_code = pthread_create(&threads[i], NULL, perform_work, &thread_args[i]);
        assert(!result_code);
    }
    printf("IN MAIN: All threads are created.\n");

//wait for each thread to complete
    for (i = 0; i < NUM_THREADS; i++) {
        result_code = pthread_join(threads[i], NULL);
        assert(!result_code);
        printf("IN MAIN: Thread %d has ended.\n", i);
    }
    printf("MAIN program has ended.\n");
    return 0;
}
```

Multi-threaded process with integer passing (output)

```
IN MAIN: Creating thread 0.  
IN MAIN: Creating thread 1.  
THREAD 0: Started.  
THREAD 1: Started.  
IN MAIN: Creating thread 2.  
THREAD 1: Will be sleeping for 5 seconds.  
THREAD 0: Will be sleeping for 3 seconds.  
THREAD 2: Started.  
THREAD 2: Will be sleeping for 4 seconds.  
IN MAIN: Creating thread 3.  
IN MAIN: Creating thread 4.  
THREAD 3: Started.  
THREAD 3: Will be sleeping for 4 seconds.  
THREAD 4: Started.  
THREAD 4: Will be sleeping for 1 seconds.  
IN MAIN: All threads are created.  
THREAD 4: Ended.  
THREAD 0: Ended.  
IN MAIN: Thread 0 has ended.  
THREAD 2: Ended.  
THREAD 3: Ended.  
THREAD 1: Ended.  
IN MAIN: Thread 1 has ended.  
IN MAIN: Thread 2 has ended.  
IN MAIN: Thread 3 has ended.  
IN MAIN: Thread 4 has ended.  
MAIN program has ended.
```

Unrestricted access to global variables!

```
#include (...)

long balance = 0; //global variable

void * inc(void * arg){
    for (long i=0; i<100000000; i++)
        balance++;
    pthread_exit(NULL);
}

void * dec(void * arg){
    for (long i=0; i<100000000; i++)
        balance--;
    pthread_exit(NULL);
}

int main(){
    pthread_t t_id1, t_id2;
    pthread_create(&t_id1, NULL, inc, NULL);
    pthread_create(&t_id2, NULL, dec, NULL);
    pthread_join(t_id1, NULL);
    pthread_join(t_id2, NULL);
    printf("Value of balance is: %ld\n", balance);
    return 0;
}
```

UNDESIRED OUTPUTS:

Value of balance is: -986105

Value of balance is: -771273

Value of balance is: 2822774

Value of balance is: 2791508

Tools to restrict access to global variables

- Mutexes: provides exclusive access to code segments/variables with the use of locks (mutual exclusion)
- Condition variables: provides synchronization and communication between threads that share a mutex

Thank you.